

University of Rome “La Sapienza”

Master in Artificial Intelligence and Robotics

Probabilistic Robotics

Prof. Giorgio Grisetti


Course web site:

http://www.dis.uniroma1.it/~grisetti/teaching/probabilistic_ro

Overview


- Probabilistic Dynamic Systems
- Dynamic Bayesian Networks (DBN)
- Inference on DBN
- Recursive Bayes Equation

Dynamic System, Deterministic View



- $f(x_{t-1}, u_{t-1})$: transition function
- $h(x_t)$: observation function
- x_{t-1} : previous state
- x_t : current state
- u_{t-1} : previous control/action
- z_t : current observation
- Δt : delay

Dynamic System, Probabilistic View



- $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{u}_{t-1})$: transition model
- $p(\mathbf{z}_t | \mathbf{x}_t)$: observation model
- \mathbf{x}_{t-1} : previous state
- \mathbf{x}_t : current state
- \mathbf{u}_{t-1} : previous control/action
- \mathbf{z}_t : current observation
- Δt : delay

Evolution of a Dynamic System




Let's start from a known initial state distribution $p(\mathbf{x}_0)$.

Evolution of a Dynamic System




A control u_0 becomes available

Evolution of a Dynamic System




The transition model $p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1})$ correlates the current states \mathbf{x}_1 with the previous control \mathbf{u}_0 and the previous state \mathbf{x}_0 .

Evolution of a Dynamic System




The observation model $p(\mathbf{z}_t \mid \mathbf{x}_t)$ correlates the observation \mathbf{z}_1 and the current state \mathbf{x}_1 .

Evolution of a Dynamic System



This leads to a recurrent structure, that depends on the *time*.

Dynamic Bayesian Networks (DBN)



- Graphical representations of stochastic dynamic processes
- Characterized by a recurrent structure

States in a DBN

The domain of the states \mathbf{x}_t , the controls \mathbf{u}_t and the observations \mathbf{z}_t are not restricted to be boolean or discrete. Examples:

- Robot localization, with laser range finder
 - states $\mathbf{x}_t \in SE(2)$, isometries on a plane
 - observations $\mathbf{z}_t \in \mathbb{R}^{\#\text{beams}}$, laser range measurements
 - controls $\mathbf{u}_t \in \mathbb{R}^2$, translational and rotational speed
- HMM
 - states $\mathbf{x}_t \in [X_1, \dots, X_{N_x}]$, finite states
 - observations $\mathbf{z}_t \in [Z_1, \dots, Z_{N_z}]$, finite observations
 - controls $\mathbf{u}_t \in [U_1, \dots, U_{N_u}]$, finite observations

Inference in a DBN requires to design a data structure that can represent a *distribution* over states.

Typical Inferences in a DBN

In a dynamic system, usually¹ we know:

- the observations made by the system $\mathbf{z}_{1:T}$, because we *measure* them
- the controls $\mathbf{u}_{1:T-1}$ because we *issue* them

Typical inferences in a DBN:

name	query	known
Filtering	$p(\mathbf{x}_T \mathbf{u}_{0:T-1}, \mathbf{z}_{1:T})$	$\mathbf{u}_{0:T-1}, \mathbf{z}_{1:t}$
Smoothing	$p(\mathbf{x}_k \mathbf{u}_{0:T-1}, \mathbf{z}_{1:T}), 0 < k < T$	$\mathbf{u}_{0:T-1}, \mathbf{z}_{1:t}$
Max a Posteriori	$\operatorname{argmax}_{\mathbf{x}_{0:T}} p(\mathbf{x}_{0:T} \mathbf{u}_{0:T-1}, \mathbf{z}_{1:T})$	$\mathbf{u}_{0:T-1}, \mathbf{z}_{1:T}$

¹usually does not mean always


Typical Inferences in a DBN

Using the traditional tools for Bayes Networks is not a good idea.

- too many variables (potentially infinite) render the solution intractable
- the domains are not necessarily discrete

However, we can exploit the recurrent structure to design procedures that take advantage of it.

DBN Inference: Filtering




Given

- the sequence of all observations $\mathbf{z}_{1:T}$ up to the current time T ,
- the sequence of all controls $\mathbf{u}_{0:T-1}$

we want to compute the distrubution over the current states

$$p(\mathbf{x}_T | \mathbf{u}_{0:T-1}, \mathbf{z}_{1:T}).$$

DBN Inference: Smoothing




Given

- the sequence of all observations $\mathbf{z}_{1:T}$ up to the current time T ,
- the sequence of all controls $\mathbf{u}_{0:T-1}$

we want to compute the distribution over a past state $p(\mathbf{x}_k | \mathbf{u}_{0:T-1}, \mathbf{z}_{1:T})$. Knowing also the controls $\mathbf{u}_{k:T-1}$ and the observations $\mathbf{z}_{k:T}$ after time k , leads to more accurate estimates than pure filtering.

DBN Inference: Maximum a Posteriori



Given

- the sequence of all observations $\mathbf{z}_{1:T}$ up to the current time T ,
- the sequence of all controls $\mathbf{u}_{0:T-1}$

we want to find the most likely *trajectory* of states $\mathbf{x}_{0:T}$.

In this case we are not seeking for a distribution. Just the most likely sequence.

DBN inference: Belief

- Algorithms for performing inference on a DBN keep track of the *estimate* of a distribution of states.
- This distribution should be stored in an appropriate data structure.
- The structure depends on
 - the knowledge of the characteristics of the distribution (e.g. Gaussian)
 - the domain of the state variables (e.g. continuous vs discrete)

When we write $b(\mathbf{x}_t)$ we mean our current belief of $p(\mathbf{x}_t|\dots)$.

The algorithms for performing inference on a DBN work by updating a belief.

DBN inference: Belief

- In the simple case of a system with discrete state $\mathbf{x} \in \{X_{1:n}\}$, the belief can be represented through an array \mathbf{x} of float values. Each cell of the array $\mathbf{x}[i] = p(\mathbf{x} = X_i)$ contains the probability of that state.
- If our system has a continuous state and we know it is distributed according to a Gaussian, we can represent the belief through its parameters (mean and covariance matrix).
- If the state is continuous but the distribution is unknown, we can use some approximate representation (e.g. weighed samples of state values).

Filtering: Bayes Recursion

We want to compute $p(\mathbf{x}_T \mid \mathbf{u}_{0:T-1}, \mathbf{z}_{1:T}) = ?$

We know:

- the observations $\mathbf{z}_{1:T}$
- the controls $\mathbf{u}_{0:T-1}$
- $p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1})$: the transition model. It is a function that given the previous state \mathbf{x}_{t-1} and control \mathbf{u}_{t-1} , tells us how likely it is to land in state \mathbf{x}_t .
- $p(\mathbf{z}_t \mid \mathbf{x}_t)$: the transition model. It is a function, that given the current state \mathbf{x}_{t-1} , tells us how likely it is to observe \mathbf{z}_t .
- $b(\mathbf{x}_{t-1})$, which is our previous belief about $p(\mathbf{x}_{t-1} \mid \mathbf{u}_{0:t-2}, \mathbf{z}_{1:t-1})$

Filtering (1)

$$p(\mathbf{x}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t}) = \quad (1)$$

splitting \mathbf{z}_t :

$$= p(\underbrace{\mathbf{x}_t}_A \mid \underbrace{\mathbf{z}_t}_B, \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C) \quad (2)$$

recall the conditional bayes rule $p(A|B, C) = \frac{p(B|A, C)p(A|C)}{p(B|C)}$

$$= \frac{p(\underbrace{\mathbf{z}_t}_B \mid \underbrace{\mathbf{x}_t}_A, \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C) p(\underbrace{\mathbf{x}_t}_A \mid \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C)}{p(\underbrace{\mathbf{z}_t}_B \mid \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C)} \quad (3)$$

Filtering: Denominator

- Let the denominator

$$\eta_t = 1/p(\mathbf{z}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}). \quad (4)$$


- Note that η_t does not depend on \mathbf{x} , thus to the extent of our computation is just a normalizing constant.
- We will come back to the denominator later.

Filtering: Observation model

- Our filtering equation becomes

$$\eta_t p(\mathbf{z}_t \mid \mathbf{x}_t, \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) p(\mathbf{x}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) \quad (5)$$

- Note that $p(\mathbf{z}_t \mid \mathbf{x}_t, \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1})$ means this



If we know \mathbf{x}_t , we do not need to know $\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}$ to predict \mathbf{z}_t , since the state \mathbf{x}_t encodes all the knowledge about the past (Markov assumption) :

$$p(\mathbf{z}_t \mid \mathbf{x}_t, \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) = p(\mathbf{z}_t \mid \mathbf{x}_t) \quad (6)$$

Filtering: Transition Model


- Thus, our current equation is

$$p(\mathbf{x}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t}) = \eta_t p(\mathbf{z}_t \mid \mathbf{x}_t) p(\mathbf{x}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) \quad (7)$$

- Still the second part of the equation is obscure.
- Our task is to manipulate it to get something that matches our preconditions.

Filtering: Transition Model

- If we would know \mathbf{x}_{t-1} , our life would be much easier, as we could repeat the trick done for the observation model:



- thus

$$p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) = p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) \quad (8)$$

Filtering: Transition Model

The sad truth is that we do not have \mathbf{x}_{t-1} , however

- Recall the following probability identities

$$p(A|C) = \sum_b p(A, B|C) \quad (9)$$

$$p(A, B|C) = p(A|B, C)p(B|C) \quad (10)$$

- combining the two above we have

$$p(A|C) = \sum_b p(A|B, C)p(B|C) \quad (11)$$

Filtering: Transition Model

The sad truth is that we do not have \mathbf{x}_{t-1} , however

- let's look again at our problematic equation, and put some letters

$$p(\underbrace{\mathbf{x}_t}_A \mid \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C) = \quad (12)$$

$$= \sum_{\mathbf{x}_{t-1}} p(\underbrace{\mathbf{x}_t}_A \mid \underbrace{\mathbf{x}_{t-1}}_B, \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C) p(\underbrace{\mathbf{x}_{t-1}}_B \mid \underbrace{\mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}}_C) \quad (13)$$

- putting in the result of Eq. 8, we highlight the transition model

$$= \sum_{\mathbf{x}_{t-1}} p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) p(\mathbf{x}_{t-1} \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) \quad (14)$$

Filtering: Wrapup

- After our efforts, we figure out that the recursive filtering equation is the following:

$$p(\mathbf{x}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t}) = \quad (15)$$

$$\eta_t p(\mathbf{z}_t \mid \mathbf{x}_t) \sum_{\mathbf{x}_{t-1}} p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) p(\mathbf{x}_{t-1} \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) \quad (16)$$

- yet, if in the last term of the product in the summation, we would not have a dependancy from \mathbf{u}_{t-1} , we would have a recursive equation.
- luckily, $p(\mathbf{x}_{t-1} \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t-1}) = p(\mathbf{x}_{t-1} \mid \mathbf{u}_{0:t-2}, \mathbf{z}_{1:t-1})$, since the last control has no influence on \mathbf{x}_{t-1} if we don't know \mathbf{x}_t .

Filtering: Wrapup

- We can finally write the recursive equation of filtering as:

$$\overbrace{p(\mathbf{x}_t \mid \mathbf{u}_{0:t-1}, \mathbf{z}_{1:t})}^{b(\mathbf{x}_t)} \quad (17)$$

$$= \eta_t p(\mathbf{z}_t \mid \mathbf{x}_t) \sum_{\mathbf{x}_{t-1}} p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) \underbrace{p(\mathbf{x}_{t-1} \mid \mathbf{u}_{0:t-2}, \mathbf{z}_{1:t-1})}_{b(\mathbf{x}_{t-1})}$$

- during the estimation, we do not have the true distribution, but rather beliefs estimate.
- Equation 18 tells us how to update a current belief once new observations/controls become available

$$b(\mathbf{x}_t) = \eta_t p(\mathbf{z}_t \mid \mathbf{x}_t) \sum_{\mathbf{x}_{t-1}} p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) b(\mathbf{x}_{t-1}) \quad (18)$$

Normalizer: η_t

- η_t is just a constant ensuring that $b(\mathbf{x}_t)$ is still a probability distribution.

$$\eta_t = \frac{1}{\sum_{\mathbf{x}_t} p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) b(\mathbf{x}_{t-1})}$$

Filtering: Discrete case

```

float transition_model(int to, int from, int control);
float observation_model(int observation, int state);

void filter(float* b, int n_states, int z, int u) {
    // clear the states
    float b_pred[n_states];
    for(int i=0; i<n_states; i++)
        b_pred[i]=0;

    // predict
    for(int i=0; i<n_states; i++)
        for(int j=0; j<n_states; j++)
            b_pred[i]+=transition_model(i,j,u)*b[j];

    // integrate the observation
    float inverse_eta = 0;
    for (int i=0; i<n_states; i++)
        inverse_eta += b_pred[i]*observation_model(z,i);

    // normalize
    float eta = 1./inverse_eta;
    for (int i=0; i<n_states; i++)
        b[i] = b_pred[i] * eta;
}

```

Filtering: Alternative Formulation

Predict: incorporate in the last belief b_{t-1} , the most recent observation.

- From transition model and the last state, compute the following joint distribution through *chain rule*

$$p(\mathbf{x}_t, \mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1}) = p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{u}_{t-1}) \underbrace{p(\mathbf{x}_t, \mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-2})}_{b_{t-1}} \quad (19)$$

- From the joint, remove \mathbf{x}_{t-1} through *marginalization*

$$\underbrace{p(\mathbf{x}_t | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1})}_{b_{t|t-1}} = \sum_{\mathbf{x}_{t-1}} p(\mathbf{x}_t, \mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1}) \quad (20)$$

Filtering: Alternative Formulation

Update: from the predicted belief $b_{t|t-1}$, compute the joint distribution that predicts the observation.

- From transition model and the last state, compute the following joint distribution through *chain rule*

$$p(\mathbf{x}_t, \mathbf{z}_t | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1}) = p(\mathbf{z}_t | \mathbf{x}_t) p(\mathbf{x}_t | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1}) \quad (21)$$

- Incorporate the current observation through *conditioning*, on the actual measurement

$$\underbrace{p(\mathbf{x}_t | \mathbf{z}_{1:t}, \mathbf{u}_{1:t-1})}_{b_{t|t}} = \frac{p(\mathbf{x}_t, \mathbf{z}_t | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1})}{p(\mathbf{z}_t | \mathbf{z}_{1:t-1}, \mathbf{u}_{1:t-1})} \quad (22)$$

Note: since we already know the value of \mathbf{z}_t , we do not need to compute the joint distribution for all possible values of $\mathbf{z} \in \mathcal{Z}$, but just for the current measurement.